

# CSC12107 – Information Systems for Business Intelligence

## Chapter 4 ETL PROCESS



KHOA CÔNG NGHỆ THÔNG TIN  
TRƯỜNG ĐẠI HỌC KHOA HỌC TỰ NHIÊN

- After complete this chapter, student will:
  - Define the ETL approach and architecture
  - Able to extract, transform and load data into a datawarehouse



# Main topic

- **Introduction to ETL**
- ETL approach and architecture
- Extracting data
- Populating the DW



# Discussion

- WHY WE HAND CODE
- WHY WE USE TOOLS
- *Are ETL systems only used to load data into the data warehouse?*




# Introduction to ETL

- ETL is the process of retrieving and transforming data from the source system and putting it into the data warehouse.
- The most underestimated and time-consuming process in DW development [?]
  - Often, 80% of development time is spent on ETL
- The ETL system is the foundation of the DW/BI project → its success makes or breaks the data warehouse.




# Introduction to ETL

- ETL stands for Extract, Transform, and Load
  - **E**xtract
    - Extract relevant data from source system
  - **T**ransform
    - Transform data to DW format; deduce new data values, validate data checks
    - Data cleansing
    - ...
  - **L**oad
    - loads the data into the DW



# ETL - Fundamental principles

- Not to slow the source system down too much
- Be careful not to disturb the source system too much
- **The extraction to be as fast as possible:**
  - **Time:** such as five minutes if we can, not three hours
  - **Size:** as small as possible, such as 10MB per day, not 1GB per day.
  - **Frequency:** once a day if we can, not every five minutes



# ETL - Fundamental principles

- The change in the source systems to be as minimal as possible
  - Should not creating triggers for capturing data changes in every single table.
- Should not have any **leakage**
- Can **recover** without data loss or damage





# Main topic

- Introduction to ETL
- **ETL approach and architecture**
- Extracting data
- Populating the DW



# ETL approach and architecture

- **Traditional approach:**

- To stage on disks or do transformation in memory

- **Alternative approaches based on:**

1. where to perform the transformations
2. who moves the data out of the source system
3. Where to put ETL processes



# ETL approach and architecture

- **Traditional approach:**



- Source system → ETL server (stage) → DW
- Source system → ETL server (no staging, in memory) → DW
- The staging area is a physical database or files.
- Putting the data into the staging area means inserting it into the database or writing it in files.



# ETL approach and architecture

## Traditional approach:

- Transforming the data in memory is faster than putting it on disk first.
- Small data?
- Big data?





# ETL approach and architecture

- **Traditional approach:**
  - To stage on disks or do transformation in memory
- **Alternative approaches based on:**
  1. **where to perform the transformations**
  2. who moves the data out of the source system
  3. Where to put ETL processes




# ETL approach and architecture

- **E-L-T:** (Extract – Load – Transform): copy the source system (OLTP) data into the data warehouse and transform it there:
  - Pull data from source system
  - Load it into DW
  - Apply the transformation by updating data in DW




4.0

# ETL approach and architecture



**Strong ETL server**  
**Strong software**



**strong DW database  
system - Usually Massively  
parallel processing (MPP)**



# ETL approach and architecture

- **Traditional approach:**
  - To stage on disks or do transformation in memory
- **Alternative approaches based on:**
  1. where to perform the transformations
  2. **who moves the data out of the source system**
  3. Where to put ETL processes









# ETL approach and architecture

1. Pulls the data out by *querying* the source system database regularly
2. Implement a *trigger* in the source system database to push the data changes out
3. Install a *schedule* process within the source system to extract data periodically.
4. Implement a procedure to *read* the database *logfiles* of the DB source to discover every data change of the DB



# ETL approach and architecture

- 4 method to execute the ETL process based on who moves the data out of the source system





# ETL approach and architecture

- **Traditional approach:**
  - To stage on disks or do transformation in memory
- **Alternative approaches based on:**
  1. where to perform the transformations
  2. who moves the data out of the source system
  3. **Where to put ETL processes**




# ETL approach and architecture

- **Three approaches based on where to put ETL processes:**
  - Execute the ETL processes in a separate ETL server that sits between the source system and the data warehouse server
  - Execute the ETL processes in the data warehouse server
  - Execute the ETL processes in the server that hosts the source system



# ETL approach and architecture

- Three approaches based on where to put ETL processes:





# Main topic

- Introduction to ETL
- ETL approach and architecture
- **Extracting data**
- Populating the DW



# Extracting database

**1. Flat file**

2. Relational database

3. Others



# Extraction from Flat files

- Example of flat file:

```
001 | Nguyễn Ngọc Thảo | Accounting  
002 | Trần Thanh Toàn | Admin  
003 | Lê Thanh Nhi | Technical support
```

- Use bulk-insert in SQL command to load data from flat file
- Quick I/O, provide the best performance
  - Bulk insert table1 from 'file1' with (field terminator = '|')





# Extraction from Flat files

- **Some remarks:**

- Must know the flat file structure:
  - Field name
  - Delimiter
  - Type and length are not fixed
  - File name
- Have access to the agreed-upon location (permission to delete/read files)
- ....



# Extracting database

- 1. Flat file**
2. Relational database
3. Others



# Extracting relational database

- **Three methods:**

- 1. Incremental extract**
2. Fixed range
3. Whole table every time



# Incremental extract

- Download only the changed rows from the source system, not the whole table
    - New row
    - Newly deleted row
    - Newly updated row
- ☐ **Based on**: **timestamp** columns, identity columns, transaction dates, **triggers**, or a combination of them



# Incremental extract

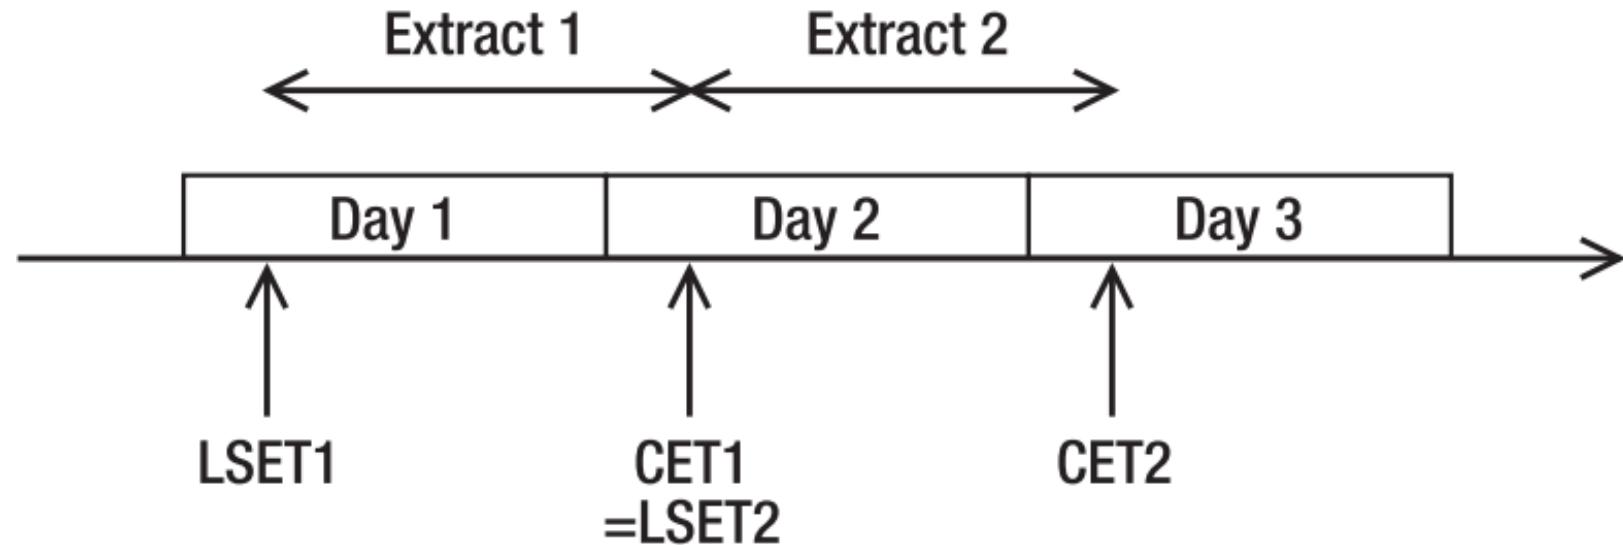
- Using a “created”/last timestamp column, updated/order date/incremental orderID
  - check whether the timestamp columns are reliable
    - contain dummy values such as 1900-01-01, Blank, null, Last updated < created date....
- Every time the row in the table changes (insert/update), the timestamp is updated

Order ID	Order Date	Some Columns	Order Status	Created	Last Updated
45433	10/10/2007	Some Data	Dispatched	10/11/2007 10:05:44	10/12/2007 11:23:41
45434	10/15/2007	Some Data	Open	10/16/2007 14:10:00	10/17/2007 15:29:02
45435	10/16/2007	Some Data	Canceled	10/16/2007 11:23:55	10/17/2007 16:19:03
...					



# Incremental extract

- Incremental extraction logic using **LSET** and **CET**
  - LSET: the time when data was last extracted.
  - CET is the time the **ETL package** started, **not** when the current **task** started





# Incremental extract

## Procedure to extract

1. Retrieve the LSET from the metadata database
2. Get the CET, which is passed in by the top-level ETL package
3. Extract the data
  1. **select** \* **from** order\_header **where** (created >= LSET and created < CET) or (last\_updated >= LSET and last\_update < CET)
4. Update meta data: writing CET as the new LSET value.



# Incremental extract

- Business rule:
  - the order happened last week but just entered into the system today (pastdated orders).
  - If we apply the previous logic to the order\_datecolumn, we will miss past-dated orders.
  - if you try to put the order date as 29 days ago → generate an error message.
  - ***select** \* **from** order\_header **where** order\_date >= (LSET – 28 days) and created < CET*





# Incremental extract

**Another way of doing incremental extract is to use the order ID:**

- Retrieve the Last Successfully Extracted ID (LSEI) from the metadata database.
- Select `max(order_id)` from `order_header`
- Set `CEI = max`
- Select `*` from `order_header` where `order_id >= LSEI` and `order_id < CEI`.
- Update meta data: `LSEI = CEI`.



# Remarks

- Incremental Extract logic is fault tolerant
- if ETL doesn't run or it failed to run, the return result will:
  - no risk of missing the data
  - Not loading data that we loaded earlier
- **How about deletion?**
  - How do we know which orders have been deleted?



# Discussion

- How about deletion?
  - soft delete, don't physically delete the record in the table
  - physically deleted
    - comparing the PK between the source table and the warehouse table
  - Using deletion trigger
    - A trigger is the most reliable approach in ETL
    - Can create separate triggers for delete, update, and insert
    - **Drawback:**



## Fixed range

- Periodically extract a certain number of records or a certain period of time based on business constraint.
  - no reliable incremental identity column
  - no timestamp columns
  - timestamp columns are not reliable
  - .....
- EX: by the end of each month, extract the data of that month



# Whole table

When:

- Deals with tables of small size
- No timestamp or identity column
- Neither incremental attribute
- No business constraint



# Extracting database

1. Flat file
2. Relational database
- 3. Others**



# Extracting Other Source Types

- XML
- spreadsheet files (Excel)
- Web logs
- Binary file
- Webservice
- Emails
- .....



# Main topic

- Introduction to ETL
- ETL approach and architecture
- Extracting data
- **Populating the DW**





# Populating the DW

- ☐ Loading the stage
- ☐ Creating the data firewall
- ☐ Populating a normalized data store (NDS)
- ☐ Populating dimension data store (DDS)
  - ☐ Populating dimension tables
  - ☐ Populating fact tables



# Stage loading (source → stage → NDS → DDS)

- Load the source system data into the stage
- Extract the data as soon as possible without doing too much transformation
  - the structure of the stage tables is similar to the source system tables
- It is better not to put any indexes or constraints in the stage database
  - to capture and report the “bad data” in the data quality process




# Stage loading

- Three different approaches of how stage tables are structured:
  1. Keeps the previous day's data in the same table
  2. Keeps each day in a separate table
  3. Uses Just one table and truncate the table every time before loading




# Stage loading

- Three different approaches of how stage tables are structured:



Approach 1



Approach 2



*don't keep the previous day's data in the stage database*

Approach 3



# Data firewall

- The data firewall is a program that checks the incoming data, similar to the firewall concept in networking → ensures data quality
  - Physically, it is an SSIS package or a stored procedure
  - Place a data firewall between the stage and the
    - **reject** the data (not load it into the DW),
    - **allow** the data (load it into the DW)
    - **fix** the data (correct the data before loading it into the DW)



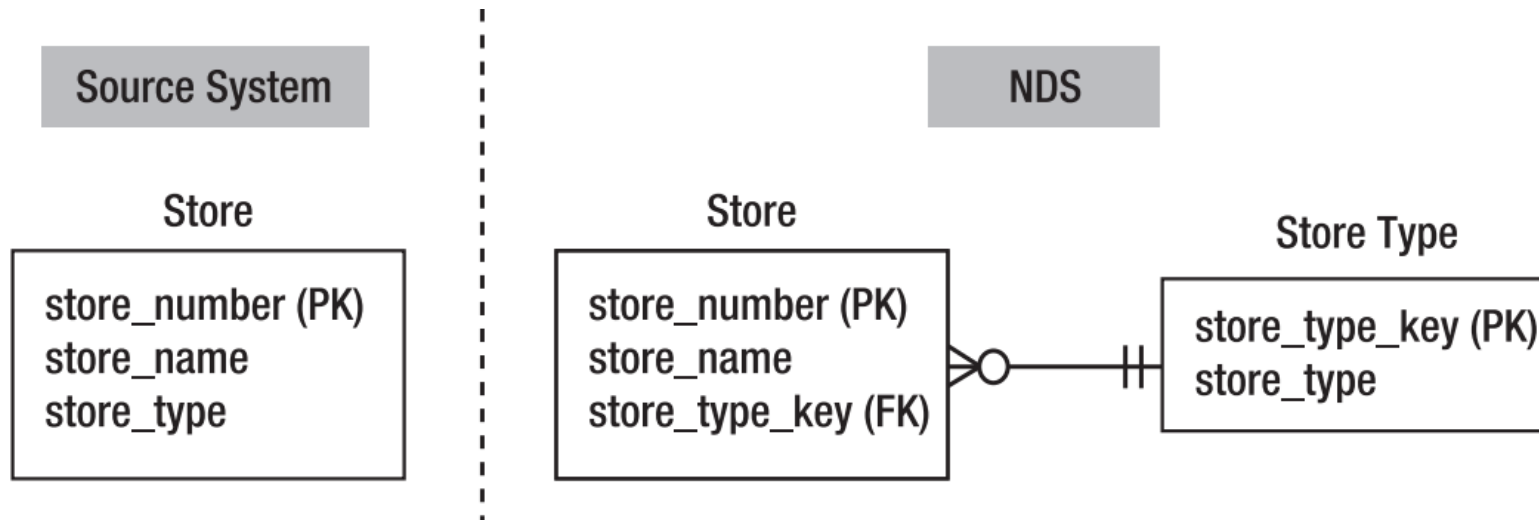
# Populating NDS (source → stage → **NDS** → DDS)

- Extract data then load it into the NDS database:
  - From the stage table
  - From the source system
- **Some remarks**
  - In the NDS, the tables are normalized
  - External data
  - Data conversion
  - Key management
  - Consider insert/ update issues



# Normalization

- Normalization is a process of removing data redundancy by implementing normalization rules (1NF, 2NF, 3 NF, BCF...)
- A normalized data store is usually in third normal form or higher





## Store

store_number	store_name	store_type
1805	Perth	Online
3409	Frankfurt	Full Outlet
1014	Strasbourg	Mini Outlet
2236	Leeds	Full Outlet
1808	Los Angeles	Full Outlet
2903	Delhi	Online

Normalized



## N\_Store

store_number	store_name	store_type_key
1805	Perth	1
3409	Frankfurt	2
1014	Strasbourg	3
2236	Leeds	2
1808	Los Angeles	2
2903	Delhi	1

## Store\_Type

store_type_key	store_type
1	Online
2	Full Outlet
3	Mini Outlet

Each store has one store type

Each store type relates to zero or many stores





# Normalization

- Data from source

store_number	store_name	store_type
2009	Dallas	Online
2237	London	Full Outlet
2014	San Francisco	Distribution Center

Store\_NDS

store_number	store_name	store_type_key
1805	Perth	1
3409	Frankfurt	2
1014	Strasbourg	3
2009	Dallas	1
2237	London	2
2014	San Francisco	4

Store\_NDS

store_number	store_name	store_type_key
1805	Perth	1
3409	Frankfurt	2
1014	Strasbourg	3

StoreType\_NDS

store_type_key	store_type
1	Online
2	Full Outlet
3	Mini Outlet

not exists in



store_type_key	store_type
1	Online
2	Full Outlet
3	Mini Outlet
4	Distribution Center

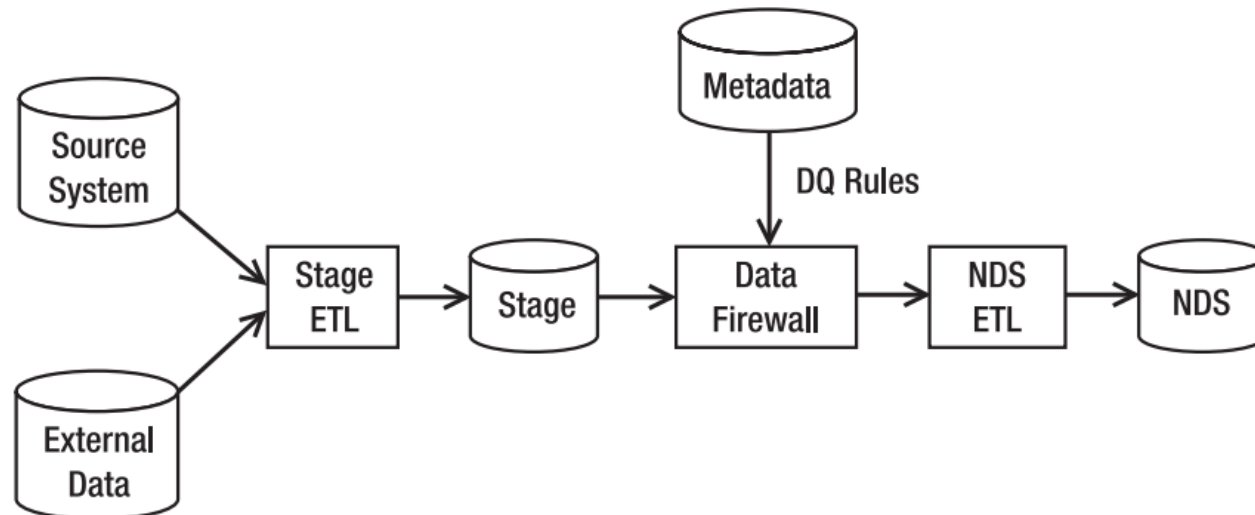




4.0

# External data

- Customer from Source system: Congo (Zaire)
  - Customer from external data: Congo
- ➔ should we **create** “Congo (Zaire)” as a new row?  
Or **replace** “Congo (Zaire)” with “Congo”?





## External data

- **TIPS:** to prevent the ETL from creating duplicate entries in the NDS country table
  - Create a rule in the data quality routine to replace “Congo (Zaire)”
  - Then the NDS ETL looks up “Congo” in the NDS country table



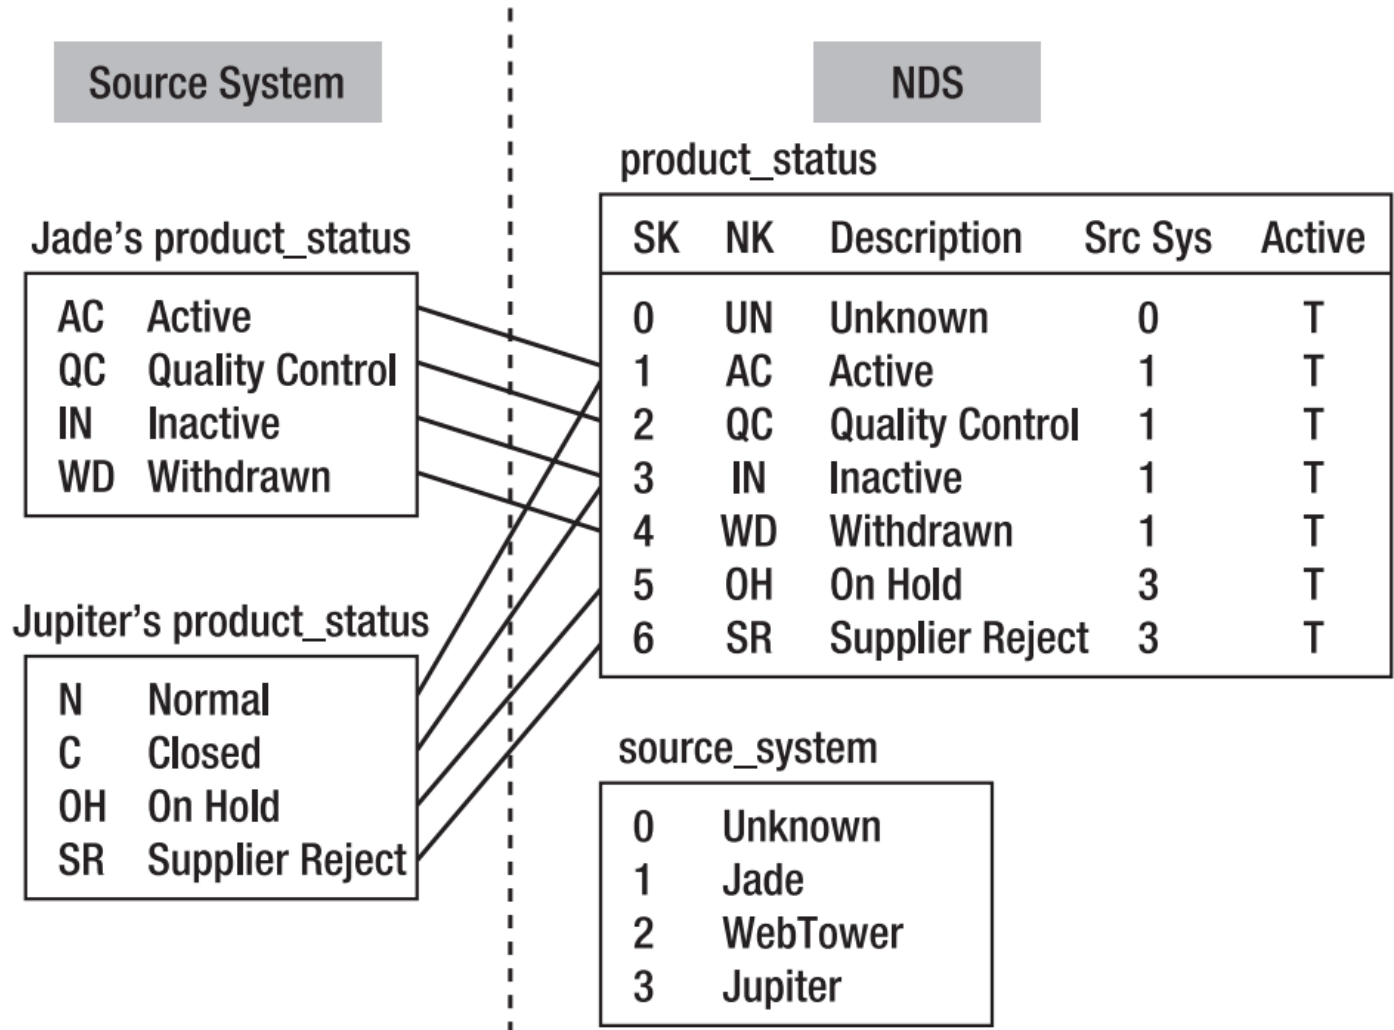
# Key management

- To be able to adapt the key changes in the source system(s).
- The data warehouse key is known as a surrogate key (SK) → enables the integration of several source systems
- The source system key is known as a natural key (NK)



# Key management

- The mapping for the product\_status table between the source systems and the NDS





# Populating DDS - **Dimension** Tables

- DDS tables: **fact** tables and **dimension** tables
- The dimension tables in the DDS are ***denormalized***
- Like populating NDS, do an **UPSERT** operation to update or insert the source row depending on whether it exists in the target
- ***Several issues need to be considered:***
  - Incremental loading,
  - Key management,
  - Denormalization,
  - Slowly changing dimension (SCD)



# Populating DDS - **Dimension** Tables

- The dimension tables in the DDS are ***denormalized***
- Like populating NDS, do an **UPSERT** operation to update or insert the source row depending on whether it exists in the target
- ***Several issues need to be considered:***
  - **Incremental loading,**
  - Key management,
  - Denormalization,
  - Slowly changing dimension

In the NDS everything is timestamped → load only the NDS rows that changed since the last ETL run



# Populating DDS - **Dimension** Tables

- The dimension tables in the DDS are ***denormalized***
- Like populating NDS, do an **UPSERT** operation to update or insert the source row depending on whether it exists in the target
- ***Several issues need to be considered:***
  - Incremental loading,
  - **Key management,**
  - Denormalization,
  - Slowly changing dimension

The Surrogate key (SK) are managed in the NDS  
→ can guarantee that all DDSs use the same SK.





# Populating DDS - **Dimension** Tables

- The dimension tables in the DDS are ***denormalized***
- Like populating NDS, do an **UPSERT** operation to update or insert the source row depending on whether it exists in the target
- ***Several issues need to be considered:***
  - Incremental loading,
  - Key management,
  - **Denormalization,**
  - Slowly changing dimension

To load the store dimension in the DDS, in the NDS we need to join the storetable with the store\_type table



# Populating DDS - **Dimension** Tables

- The dimension tables in the DDS are ***denormalized***
- Like populating NDS, do an **UPSERT** operation to update or insert the source row depending on whether it exists in the target
- ***Several issues need to be considered:***
  - Incremental loading,
  - Key management,
  - Denormalization,
  - **Slowly changing dimension (SCD)**

SCD type 1(overwrite),  
SCD type 2 (rows),  
SCD type 3 (column)



# Populating DDS - **Fact** Tables

- Fact tables are normally large tables
- Table partitioning can speed up the update operation significantly when correctly applied to a fact table
- The dimensional key index is required to find the row that we want to update



# Assuring Data Quality

- when building a DW, it is important to think about data quality **as early as possible**
- The data quality process includes the activities to make sure the data in the DW is correct and complete
  - ☐ set up rules that define what “bad data”
  - ☐ Reporting
  - ☐ Monitoring
  - ☐ Cleaning/ Correcting



# Assuring Data Quality

- Example: customers can purchase a product, or they can subscribe to a package
  - first subscription date > last cancellation date → invalid condition
  - Question:
    - The last cancellation date is wrong
    - The first subscription date is wrong
    - or both??



# Assuring Data Quality

- ☐ Data Cleansing and Matching
- ☐ Cross-checking with External Sources
- ☐ Data Quality Rules



# Data Cleansing and Matching

- **Data cleansing**, or **data scrubbing** is the process of identifying and correcting **dirty data**
  - Dirty data means incomplete, wrong, duplicate, or out-of-date data
- **Ex:**
  - Checking stores' tables to make sure the store names, store numbers, store types, and store addresses are all correct
  - Making sure that there are no duplicate customer records → **data matching**
    - use = for matching numeric data. For example, "if A = B, then...."
    - is 5.029 the same as 5.03?
    - 03/01/2008 the same as 01/03/2008?



# SQL server matching

- Three types of matching logic:
  - **Exact**: all characters are the same, for example “Los Angeles” and “Los Angeles.”
  - → using a **Lookup** transformation
  - **fuzzy** (approximate): finds how similar a set of data is to another set of data
    - using the Fuzzy Lookup
    - “You can’t hurry love” and “You cannot hurry love” have a similarity score of 0.81666672 and a confidence level of 0.59414238
  - **rule based**: use certain rules and data to identify a match
    - In product names “movie” is the same as “film”. → implemented with Script Component
    - “For product code, omit the spaces when comparing” so that “KL 7923 M” is the same as “KL7923M.”